



Deep learning on ML Engine

Balázs Tóth

Software Engineer at Trustpilot

About me

- Software engineer at Trustpilot
- Applied machine learning to detect fake reviews and fraudulent user behavior

Agenda

Overview of an ML engine service in the fraud detection context

A tutorial on training neural networks with Keras on ML Engine

The problem

- Limitations of your local computer
- Limitations of your time

Why do we use ML Engine in Trust and Transparency?

- Train multiple models in parallel
- Easy access to powerful computers
- Easy to deploy and maintain models
- Hyperparameter optimization
- Good ecosystem with Dataflow and BigQuery

Why do we use ML Engine in Trust and Transparency?

- Train multiple models in parallel
- Easy access to powerful computers
- Easy to deploy and maintain models
- Hyperparameter optimization
- Good ecosystem with Dataflow and BigQuery

Why do we use ML Engine in Trust and Transparency?

- Train multiple models in parallel
- Easy access to powerful computers
- Easy to deploy and maintain models
- Hyperparameter optimization
- Good ecosystem with Dataflow and BigQuery

Why do we use ML Engine in Trust and Transparency?

- Train multiple models in parallel
- Easy access to powerful computers
- Easy to deploy and maintain models
- Hyperparameter optimization
- Good ecosystem with Dataflow and BigQuery

Why do we use ML Engine in Trust and Transparency?


- Train multiple models in parallel
- Easy access to powerful computers
- Easy to deploy and maintain models
- Hyperparameter optimization
- Good ecosystem with Dataflow and BigQuery

Overview of an ML Engine Service

Spammer adversaries in Trustpilot



Solution: Spam review detector



Evolution of spam reviews with the iterations of the spam review detector

Machine learning workflow

- Preprocessing - Google Dataflow
- Training and evaluation - Google ML Engine
- Deployment of the model - Google ML Engine API
- Batch prediction (2nd part of evaluation) - Google ML Engine

Machine learning workflow

- Preprocessing - Google Dataflow
- Training and evaluation - Google ML Engine
- Deployment of the model - Google ML Engine API
- Batch prediction (2nd part of evaluation) - Google ML Engine

Machine learning workflow

- Preprocessing - Google Dataflow
- Training and evaluation - Google ML Engine
- Deployment of the model - Google ML Engine API
- Batch prediction (2nd part of evaluation) - Google ML Engine

Machine learning workflow

- Preprocessing - Google Dataflow
- Training and evaluation - Google ML Engine
- Deployment of the model - Google ML Engine API
- Batch prediction (2nd part of evaluation) - Google ML Engine

A tutorial on training neural networks with Keras on ML Engine

Example application: Classification of newsgroups



Public Dataset:

<http://qwone.com/~jason/20Newsgroups/>

Code:

<https://github.com/tothbalazs0920/ml-engine-example>

Use only 4 groups

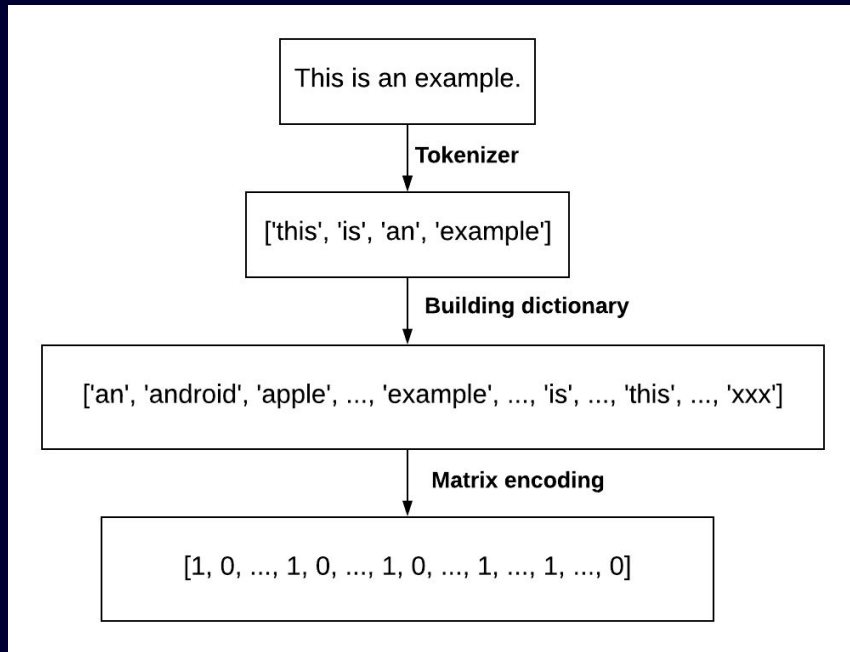
- atheism
- religion
- graphics
- science

1. Setup a project on Google Cloud Console



- Enable ML Engine API's
- Install Google cloud command line tools

2. Preprocess the data - bag of words



2. Preprocess the data and upload the preprocessed data to a gcs

```
tokenizer = keras.preprocessing.text.Tokenizer(num_words=300)
tokenizer.fit_on_texts(news_train)
encoded_docs = tokenizer.texts_to_matrix(news_train, mode='count')
```


3. Create a Keras model

```
def create_model(number_of_features):  
  
    feature_vector = Input(shape=(number_of_features,), name='feature_vector')  
  
    layer1 = Dense(10, activation='relu')(feature_vector)  
    layer2 = Dense(10, activation='relu')(layer1)  
    predictions = Dense(4, activation='softmax', name='predictions')(layer2)  
  
    model = Model(inputs=[feature_vector], outputs=[predictions])  
    model.compile(optimizer='adam', loss={'predictions': 'binary_crossentropy'},  
                  metrics=['accuracy'])  
    return model
```

4. Set up the training in Keras

- `model.fit()` and use a machine with a lot of memory (256 GB)
- `model.fit_generator()` and use a machine with less memory

4. Set up the training in Keras

- `model.fit()` and use a machine with a lot of memory (256 GB)
- `model.fit_generator()` and use a machine with less memory

5. Add arguments to the entry file

```
if __name__ == "__main__":
    parser = argparse.ArgumentParser()
    parser.add_argument('--job-dir',
                        required=True,
                        type=str,
                        help='GCS to write checkpoints')
    parser.add_argument('--x_train_file',
                        required=True,
                        type=str,
                        help='path of training samples')
    parser.add_argument('--y_train_file',
                        required=True,
                        type=str,
                        help='path of target result of training samples')
    parser.add_argument('--x_test_file',
                        required=True,
                        type=str,
                        help='path of test samples')
    parser.add_argument('--y_test_file',
                        required=True,
                        type=str,
                        help='path of target result of test samples')
```

6. Convert the keras model to tensorflow model after training

```
def save_tensorflow_model(model, export_path):  
    if file_io.file_exists(export_path):  
        return  
    builder = saved_model_builder.SavedModelBuilder(export_path)  
    signature = predict_signature_def(inputs={'input': model.inputs[0]},  
                                     outputs={'output': model.outputs[0]})  
  
    with K.get_session() as sess:  
        builder.add_meta_graph_and_variables(  
            sess=sess,  
            tags=[tag_constants.SERVING],  
            signature_def_map={  
                signature_constants.DEFAULT_SERVING_SIGNATURE_DEF_KEY: signature  
            }  
        )  
        builder.save()
```

7. Create a python package and upload it to a gcs

```
python setup.py sdist --formats=gztar
```


8. Submit a training job to ML Engine

```
gcloud ml-engine jobs submit training sentiment_analysis_1
--region europe-west1
--runtime-version 1.6
--config config.yml
--package-path trainer
--module-name trainer.task
--job-dir your-bucket-name/jobdir/
--packages your-bucket-name/package/latest.tar.gz
--
--x_train_file your-bucket-name/data/x_train.csv
--y_train_file your-bucket-name/data/y_train.csv
--x_test_file your-bucket-name/data/x_test.csv
--y_test_file your-bucket-name/data/y_test.csv
--number_of_epochs 200
--number_of_features 100
--number_of_training_examples 25000
--number_of_test_examples 25000
--batch_size 2000
--model_location sentiment_analysis_1
```

Hyperparameter tuning

```
trainingInput:
  pythonVersion: '3.5'
  region: eu-west1
  hyperparameters:
    goal: MAXIMIZE
    hyperparameterMetricTag: accuracy
    maxTrials: 4
    maxParallelTrials: 2
    params:
      - parameterName: number_of_nodes_one
        type: INTEGER
        minValue: 10
        maxValue: 100
        scaleType: UNIT_LINEAR_SCALE
      - parameterName: number_of_nodes_two
        type: INTEGER
        minValue: 10
        maxValue: 100
        scaleType: UNIT_LINEAR_SCALE
```


Interpreting the model with SHAP



Thank you

Email: bat@trustpilot.com

LinkedIn: <https://www.linkedin.com/in/balazs-toth-6186005a/>